

Amazon

Santa Clara, CA 95054

Website: <https://damon-demon.github.io>

Mobile: (+1)-347-574-5875

Email: [damonzym@amazon.com](mailto:damonzym@amazon.com)[\[Google Scholar\]](#)

---

**RESEARCH FOCUSES**

---

**Deep learning:** Generative models (LLMs, multi-modality, diffusion models), AI Safety (adversarial attack & defense, machine unlearning)**Optimization:** Sparsity learning for model/dataset compression, Black-box optimization

---

**EDUCATION**

---

**Ph.D. in Computer Science, Michigan State University**

Jan. 2021– Oct. 2025

**M.S. in Electrical Engineering, Columbia University**

Aug. 2018– Dec. 2019

**B.Eng in Electronic and Electrical Engineering, University of Sheffield**

Sep. 2015– July 2018

---

**SELECTED PUBLICATIONS**

---

[Google Scholar](#) (\* represents equal contribution)

- [1] **Y. Zhang**, T. Wang, J. Gesi, Z. Wang, Y. Lu, J. Lin, S. Zhan, V. Gao, R. Jiao, J. Liu, K. Qian, Y. Tang, R. Xue, H. Zhang, Q. Cui, Y. Guo, D. Wang, “[Shop-R1: Rewarding LLMs to Simulate Human Behavior in Online Shopping via Reinforcement Learning](#)”, Under Review.
- [2] **Y. Zhang**, T. Zhi, J. Liu, S. Sang, L. Jiang, Q. Yan, S. Liu, L. Luo, “[ID-Patch: Robust ID Association for Group Photo Personalization](#)”, *CVPR’25*
- [3] **Y. Zhang**, X. Chen, J. Jia, Y. Zhang, C. Fan, J. Liu, M. Hong, K. Ding, S. Liu, “[Defensive Unlearning with Adversarial Training for Robust Concept Erasure in Diffusion Models](#)”, *NeurIPS’24*
- [4] **Y. Zhang\***, J. Jia\*, X. Chen, A. Chen, Y. Zhang, J. Liu, K. Ding, S. Liu, “[To Generate or Not? Safety-Driven Unlearned Diffusion Models Are Still Easy To Generate Unsafe Images ... For Now](#)”, *ECCV’24*
- [5] A. Chen\*, **Y. Zhang\***, J. Jia, J. Diffenderfer, J. Liu, K. Parasyris, Y. Zhang, Z. Zhang, B. Kailkhura, S. Liu, “[DeepZero: Scaling up Zeroth-Order Optimization for Deep Model Training](#)”, *ICLR’24*
- [6] **Y. Zhang**, X. Chen, J. Jia, S. Jia, K. Ding “[Text-Visual Prompting for Efficient 2D Temporal Video Grounding](#)”, *CVPR’23*
- [7] **Y. Zhang\***, A.K. Kamath\*, Q. Wu\*, Z. Fan\*, W. Chen, Z. Wang, S. Chang, C. Hao, S. Liu, “[Data-Model-Circuit Tri-Design for Ultra-light Video Intelligence on Edge Devices](#)”, *ASP-DAC’23*
- [8] **Y. Zhang**, Y. Yao, J. Jia, J. Yi, M. Hong, S. Chang, S. Liu, “[How to Robustify Black-Box ML Models? A Zeroth-Order Optimization Perspective](#)”, International Conference on Learning Representation (*ICLR’22 - Spotlight, acceptance rate 5%*)

---

**WORK EXPERIENCE**

---

**Applied Scientist, Amazon**

Oct. 2025– Present.

---

**RESEARCH EXPERIENCE**

---

**Human Online Shopping Behavior Simulation via RL**

June. 2025 - July. 2025

Supervisor: [Dakuo Wang](#) (NEU), [Jiri Gesi](#) (Amazon)

- Introduce RL into a simulation-oriented human online shopping behavior modeling task.
- Develop a reinforcement-learning framework with a hybrid reward design. It integrates a self-certainty signal for rationale generation with a hierarchical reward scheme for action prediction.
- **Publications:** [\[1\]](#)

**Multi-ID Consistency for Personalized Diffusion Model**

May. 2024 - Nov. 2024

Supervisor: [Tiancheng Zhi](#) (ByteDance)

- Explore how to link face ID features with their corresponding locations using visual patches in conditioning images, ensuring better resemblance and accurate position control without ID leakage.
- Removal of the reliance on auxiliary segmentation models, requiring only a single point for ID position control, as opposed to segmented masks or head bounding boxes.
- **Publications:** [\[2\]](#)

**Adversarial Unlearning for Diffusion Model**

Nov. 2023 - May. 2024

Supervisor: [Sijia Liu](#) (MSU)

- Explore the integration of AT with concept erasing (or machine unlearning) in DMs.
- Design a utility-retaining regularization using curated external retain prompt data to balance the trade-off between effective unlearning and high-quality image generation.
- **Publications:** [\[3\]](#)

**Robustness Evaluation for Unlearned Diffusion Models**

May. 2023 - Oct. 2023

Supervisor: [Sijia Liu](#) (MSU), [Xin Chen](#) (Intel)

- Propose an evaluation framework built upon adversarial attacks (also referred to as adversarial prompts), in order to discern the trustworthiness of these safety-driven unlearned DMs.
- Develop a novel adversarial learning approach called UnlearnDiff that leverages the inherent classification capabilities of DMs to streamline the generation of adversarial prompts.
- **Publications:** [\[4\]](#)

**Scalable Model Training without Backpropogation**

Jan. 2023 - May. 2023

Supervisor: [Sijia Liu](#) (MSU)

- Propose a sparsity-induced ZO training protocol that extends the model pruning methodology using only finite differences to explore and exploit the sparse DL prior in CGE.
- **Publications:** [\[5\]](#)

**Efficient 2D Temporal Video Grounding (TVG)**

May.- Dec. 2022

Supervisor: [Xin Chen](#) (Intel)

- Propose an effective and efficient framework to train 2D TVG models, in which we leverage text-visual prompting (TVP) to improve the utility of sparse 2D visual features
- **Publications:** [\[6\]](#)

**Model Compression for Object Tracking**

Sept. 2021 - May. 2022

Supervisor: [Sijia Liu](#) (MSU)Collaborator: [Callie Hao](#)(Georgia Tech), [Shiyu Chang](#)(UCSB), [Zhangyang Wang](#)(UT Austin)

- Saliency-guided spatial data reduction method is devised to eliminate uninformative pixels from both the input frames as well as the intermediate feature maps
- Utilizing kernel-wise pattern-aware model sparsity to achieve hardware-friendly model compression.
- **Publications:** [\[7\]](#)